

Robust Gamma-KalmanNet State Estimation for QCar under Sensor Faults

Quang Huy Nugyen

Echrak Chnib

September 30, 2026

Archived GRU description and earlier results

Abstract

This article presents a robust KalmanNet-style state estimator for the QCar platform. The estimator preserves a physics-based vehicle prediction model and learns the uncertain part of the filtering problem: how much each measurement should be trusted and how strongly the state should be corrected. The learned measurement trust vector is denoted by γ_k and acts as a “robust gamma” gate. The proposed estimator tracks $\mathbf{x} = [x, y, \psi, v]^T$ using GPS position, fused heading, tachometer velocity, IMU, steering, and wheel-speed information. Training combines real QCar driving data with synthetic GPS, heading, velocity, IMU, steering, and wheel-speed faults. The learned estimator reduces attacked validation loss from 1.1095 to 0.4028 and improves four-state RMSE compared with a kinematic reference in offline validation. The result is a hybrid observer: vehicle physics provides the prior state, while a recurrent neural update learns when corrupted measurements should be down-weighted.

1 Introduction

Autonomous QCar experiments require a state estimate that is accurate enough for feedback control and stable enough during bad sensor periods. A standard Extended Kalman Filter (EKF) is a strong baseline when its assumptions are correct. It uses a vehicle model for prediction and a measurement model for correction. However, the EKF can become fragile when the measurement noise is not just random noise but a structured fault: GPS jumps, GPS dropout, frozen velocity, wheel speed scaling, heading glitches, or short IMU bursts. In those cases the filter may apply a correction in the wrong direction because the measurement is still treated as believable.

This work develops a robust KalmanNet-style estimator for the QCar. The key design choice is simple:

keep the vehicle model for prediction, and learn a trust-controlled update for difficult sensor behavior.

The learned trust vector γ_k is the important element. When a channel is clean, its trust should stay high. When a channel is attacked or unavailable, its trust should decrease so that the innovation cannot force a large state jump. This is close in spirit to KalmanNet [?] and RobustStateNet [?], but the proposed QCar estimator is more conservative: the prediction step is an analytical bicycle model rather than a learned predictor.

2 Problem Definition

2.1 State and Measurement

The estimated state is

$$\mathbf{x}_k = [x_k \quad y_k \quad \psi_k \quad v_k]^T, \quad (1)$$

where x and y are global position, ψ is heading, and v is forward velocity. Yaw rate is recorded as an auxiliary signal in some experiments, but the estimator evaluated in this article uses the four-state representation above.

The measurement vector is

$$\mathbf{z}_k = [x_{gps,k} \ y_{gps,k} \ \psi_{fused,k} \ v_{tach,k}]^T, \quad \mathbf{z}_k = H\mathbf{x}_k + \mathbf{n}_k, \quad H \approx I_4. \quad (2)$$

The measured heading is reconstructed using the same heading-fusion procedure as the onboard estimator. GPS availability metadata is retained because it is essential for distinguishing sensor faults from normal innovation.

2.2 Raw Sensor Inputs

The model receives raw inputs grouped by physical source:

- IMU branch: velocity context, heading context, acceleration, and yaw rate;
- steering branch: velocity context, heading context, and steering angle;
- wheel branch: velocity context, heading context, and four wheel-speed channels;
- availability metadata: GPS valid flag, heading valid flag, GPS hold flag, and GPS age.

This separation matters because each sensor family has different failure behavior. GPS position can jump or disappear. Velocity can freeze or become scaled. IMU signals can become noisy or biased. The learned update uses these signals to decide whether the current measurement should be trusted.

3 Estimator Architecture

Figure ?? shows the estimator at a high level. The prediction side is deterministic. The update side is learned and recurrent.




Figure 1: Hybrid RKNet structure used in this work. The prior is generated by QCar physics. The neural update learns a measurement trust gate and a Kalman-style gain.

3.1 Kinematic Prediction

The prediction step has no learned parameters. It is a bicycle-model update with wheelbase L :

$$\dot{\psi}_k = \frac{v_k \tan(\delta_k)}{L}, \quad (3)$$

where δ_k is the steering angle. The predicted pose is

$$x_k^- = x_{k-1}^+ + v_{pose,k} \cos(\psi_{k-1}^+) \Delta t, \quad (4)$$

$$y_k^- = y_{k-1}^+ + v_{pose,k} \sin(\psi_{k-1}^+) \Delta t, \quad (5)$$

$$\psi_k^- = \text{wrap} \left(\psi_{k-1}^+ + \frac{v_{pose,k} \tan(\delta_k)}{L} \Delta t \right). \quad (6)$$

Longitudinal velocity is predicted with a first-order lag model from the previous velocity and throttle:

$$\dot{v}_k = \frac{-v_{k-1}^+ + g u_k}{\tau}, \quad v_k^- = \text{clip}(v_{k-1}^+ + \dot{v}_k \Delta t). \quad (7)$$

The numerical parameters used in the experiments are $L = 0.2$, $\tau = 0.301$, $g = 6.598$, $|v| \leq 2.0$ m/s, and maximum acceleration 2.0 m/s².

3.2 Learned Robust Update

The learned update receives the prior $\mathbf{x}_k^-$, the previous corrected state $\mathbf{x}_{k-1}^+$, the current measurement $\mathbf{z}_k$, the previous measurement $\mathbf{z}_{k-1}$, and availability metadata. It forms three main residual terms:

$$\Delta \mathbf{x}_k = \text{wrap}(\mathbf{x}_k^- - \mathbf{x}_{k-1}^+), \quad (8)$$

$$\mathbf{r}_k = \text{wrap}(\mathbf{z}_k - H \mathbf{x}_k^-), \quad (9)$$

$$\Delta \mathbf{z}_k = \text{wrap}(\mathbf{z}_k - \mathbf{z}_{k-1}). \quad (10)$$

The hard availability vector is

$$\mathbf{a}_k = [g_k \quad g_k \quad h_k \quad 1]^T, \quad (11)$$

where g_k is the GPS-position valid flag and h_k is the heading valid flag. Unavailable measurement channels are removed before the learned correction.

The update feature vector is

$$\mathbf{f}_k = [\Delta \mathbf{x}_k / \mathbf{s}_x, (\mathbf{a}_k \odot \mathbf{r}_k) / \mathbf{s}_z, (\mathbf{a}_k \odot \Delta \mathbf{z}_k) / \mathbf{s}_z, \mathbf{a}_k, g_k^{hold}, \log(1 + \min(age_k, 10)), \log(1 + |\mathbf{a}_k \odot \mathbf{r}_k|), \log(1 + |\mathbf{a}_k \odot \Delta \mathbf{z}_k|)]. \quad (12)$$

For the four-state model this gives a 26-dimensional update feature vector. The scale vectors used for normalization are

$$\mathbf{s}_x = \mathbf{s}_z = [0.25, 0.25, 0.20, 0.50].$$

The robust gamma/trust vector is generated by a small MLP:

$$\boldsymbol{\gamma}_k = \sigma(\text{MLP}_\gamma(\mathbf{f}_k)). \quad (13)$$

Only the measurement part of this vector is used directly as the measurement trust gate. Clean channels should stay close to one; attacked channels should move toward zero.

The Kalman-style gain is generated by a GRU and an MLP:

$$\mathbf{h}_k = \text{GRU}(\boldsymbol{\gamma}_k \odot \mathbf{f}_k, \mathbf{h}_{k-1}), \quad (14)$$

$$K_k = s \tanh(\text{MLP}_K(\mathbf{h}_k)). \quad (15)$$

The gain output is bounded with $s = 1.2$. Exponential smoothing is also applied to the gain and trust gate with $\alpha = 0.35$.

The corrected state is

$$\mathbf{x}_k^+ = \mathbf{x}_k^- + K_k (\mathbf{a}_k \odot \boldsymbol{\gamma}_k^z \odot \mathbf{r}_k), \quad (16)$$

where $\boldsymbol{\gamma}_k^z$ is the measurement-trust part of the learned gate. Equation (??) is the main robustness mechanism. A corrupted GPS or velocity measurement may create a large innovation, but the innovation is first multiplied by the learned trust gate before it reaches the gain. The runtime also clamps the per-step correction to

$$[0.35, 0.35, 0.35, 0.80]$$

for $[x, y, \psi, v]$.

4 Dataset and Training

4.1 Experimental Data

The experimental data consists of time-aligned QCar sensor signals, measurement vectors, reference state estimates, timestamps, and sensor availability metadata. The main training sequence contains 52,921 samples over 671.84 seconds of driving, with a mean timestep of 0.0127 seconds. The supervised target is a clean EKF reference estimate rather than motion-capture ground truth.

Table 1: Main experimental data and training settings.

Item	Value
State order	$[x, y, \psi, v]$
Measurement order	$[x_{gps}, y_{gps}, \psi_{fused}, v_{tach}]$
Training sequence	Real QCar driving data
Samples / duration	52,921 samples / 671.84 s
Mean timestep	0.0127 s
Validation split	final 20% of sequence
Training windows	$T = 50$ for Phase B, $T = 80$ for Phase C
Predictor mode	kinematic QCar bicycle model
Updater hidden size	32
Gain hidden size	32
Optimizer learning rate	1.0×10^{-4}

4.2 Training Phases

Because the predictor used in this study is kinematic, predictor pretraining is not required. The useful training phases are:

- Phase B: freeze the predictor and train only the learned updater;
- Phase C: fine tune the complete recurrent runtime loop with teacher forcing disabled.

The state loss is weighted MSE:

$$L_{state} = \frac{1}{N} \sum_k \left\| W^{1/2} (\hat{\mathbf{x}}_k - \mathbf{x}_k^*) \right\|_2^2, \quad W = [1.0, 1.0, 2.5, 1.2]. \quad (17)$$

The full loss combines state accuracy, trust supervision, gain regularization, and temporal smoothness:

$$\mathcal{L} = \lambda_u L_{upd} + \lambda_p L_{pred} + \lambda_\gamma L_\gamma + \lambda_K L_K + \lambda_{\Delta K} L_{\Delta K} + \lambda_{\Delta \gamma} L_{\Delta \gamma}. \quad (18)$$

The Phase C weights are $\lambda_u = 0.85$ and $\lambda_p = 0.15$. Measurement-mask supervision uses a clean target of one and an attacked target of zero for each affected measurement channel.

5 Sensor Attack Augmentation

The training data is not only clean driving. Synthetic attacks are injected offline so the model sees examples of sensor faults and receives supervision for which channels should be trusted.

The GPS attacks are especially important for QCar because GPS position directly drives x and y correction. A jump attack can make the EKF snap to a wrong position unless the measurement is down-weighted. A dropout attack is different: the availability gate should remove x/y , while the learned update continues using the remaining channels.

Table 2: Sensor-fault augmentation configuration used in the experiments.

Attack group	Probability	Enabled failure types
Raw branch attack	0.25	bias, scale, freeze, noise, ramp, zero-out
Direct measurement attack	0.35	heading and velocity corruption
GPS position attack	0.45	noise, freeze, jump, dropout, reacquisition
Maximum raw branches attacked	1	one of IMU, steering, wheel

6 Results

6.1 Training History

The training history contains 20 logged epochs after predictor pretraining was omitted. The model selected for robustness corresponds to epoch 15 in Phase B. Table ?? summarizes the main points.

Table 3: Training history. Selection loss combines clean and attacked validation, with attacked validation weighted more strongly.

Model point	Epoch	Clean val.	Attacked val.	Selection loss
First logged epoch	1	0.00668	1.10951	0.66838
Selected robust model	15	0.01751	0.40279	0.24868
Final Phase C model	20	0.02408	0.43457	0.27037

The attacked validation loss decreases from 1.10951 to 0.40279 at the selected robust model, a reduction of about 63.7%. The clean validation loss is higher at the robust model than at the first clean-only epoch. This is an expected tradeoff: the estimator becomes more conservative so it can survive corrupted measurements.

6.2 Offline Validation

Offline validation compares a kinematic reference and the learned RKNet model against the EKF target on an independent driving sequence. Some recorded signals contain an auxiliary yaw-rate channel, so Table ?? reports only the four-state surface used by the proposed estimator.

Table 4: Offline RMSE against the EKF target on an independent validation sequence.

Method	x [m]	y [m]	ψ [rad]	v [m/s]	Mean, four states
Kinematic reference	0.0663	0.0574	0.1151	0.0134	0.0631
Learned RKNet	0.0076	0.0168	0.1096	0.0124	0.0366

The learned estimator strongly improves x and y RMSE and gives a smaller four-state mean RMSE. Heading remains the hardest channel. This is reasonable because heading is affected by angle wrapping, GPS heading availability, and the reference itself.

6.3 Runtime Comparator Run

A representative runtime comparison contains 8,413 model samples over 95.97 seconds, with no fallback samples. It includes five attack segments and 1,510 attack-active samples. Only 421 samples had fresh GPS, while 8,380 used GPS hold validity, so this run is useful for checking robustness during limited fresh GPS availability.



kalmarobust_result-12.png

Figure 2: State-estimation behavior during the runtime comparison. The robust estimator follows the clean-reference EKF while the ordinary EKF shows larger spikes during attack intervals.



kalmarobust_result-13.png

Figure 3: Gain and measurement-mask behavior during the runtime comparison. The learned mask and effective gain decrease during GPS jump and wheel-bias intervals, reducing dependence on corrupted measurements.

Figures ?? and ?? show the main behavior expected from the robust gamma gate. During the GPS jump intervals, the measurement mask for position decreases and the effective diagonal gain also drops. During wheel-bias intervals, the velocity-side trust and gain become more conservative. This means the update is not blindly following the measurement; it is using the innovation only after passing it through the learned reliability gate.

7 Discussion

The important engineering result is the separation of responsibilities. The QCar prediction step carries the vehicle kinematics and basic longitudinal behavior. The neural update does not need to relearn the full vehicle model. Instead, it learns the part that is hard to tune manually: the measurement trust policy under abnormal sensor conditions.

The learned trust mechanism also makes the estimator more interpretable than a single end-to-end recurrent network. The plotted γ_k values show when the estimator is reducing confidence in

GPS position, heading, or velocity. The gain plots show whether the final correction is aggressive or conservative. This is useful for debugging because a bad output can be traced to three quantities: the prior state, the innovation, and the trust/gain pair.

There are also limitations:

- the target is an EKF/reference estimator, not absolute ground truth;
- the attacks are synthetic, so real hardware failures still need more dedicated experiments;
- the present predictor is kinematic, so raw-branch learned feature masking is not used in the prediction path;
- the model is trained mainly on the available QCar driving sequences and should be rechecked before claiming generalization to other vehicles or tracks;
- the clean/robust tradeoff is visible: robustness improves attacked validation loss, while clean validation loss becomes slightly worse.

8 Conclusion

The robust gamma-KalmaNet estimator developed in this work is a practical hybrid state estimator for QCar. It preserves a deterministic QCar kinematic prediction model and learns a recurrent Kalman-style update with a measurement trust gate. The trust gate suppresses corrupted GPS, heading, and velocity channels before the innovation is applied to the state. Training with synthetic sensor attacks reduces attacked validation loss substantially, and offline validation shows lower four-state RMSE than the kinematic reference on an independent validation sequence. The result is an estimator that is still close to the filtering structure used in classical observers, but with a learned robustness layer where hand-tuned noise models are weakest.

Data and Code Availability

The implementation and experimental material are maintained in the project repository. The results in this article are based on real QCar driving sequences with synthetic sensor-fault augmentation. The reported metrics can be reproduced by training the robust update on the described data split and evaluating the selected model on an independent validation sequence.

References

- [1] G. Revach, N. Shlezinger, X. Ni, A. L. Escoriza, R. J. G. van Sloun, and Y. C. Eldar, “KalmanNet: Neural network aided Kalman filtering for partially known dynamics,” *IEEE Transactions on Signal Processing*, vol. 70, pp. 1532–1547, 2022.
- [2] P. Dahal, S. Mentasti, L. Paparusso, S. Arrigoni, and F. Braghin, “RobustStateNet: Robust ego vehicle state estimation for Autonomous Driving,” *Robotics and Autonomous Systems*, vol. 172, article 104585, 2024.